

# ENDIORBOT

## MASTER PLAN v3.1 FINAL

*SDLC Control Plane for Claude Code Workflow*



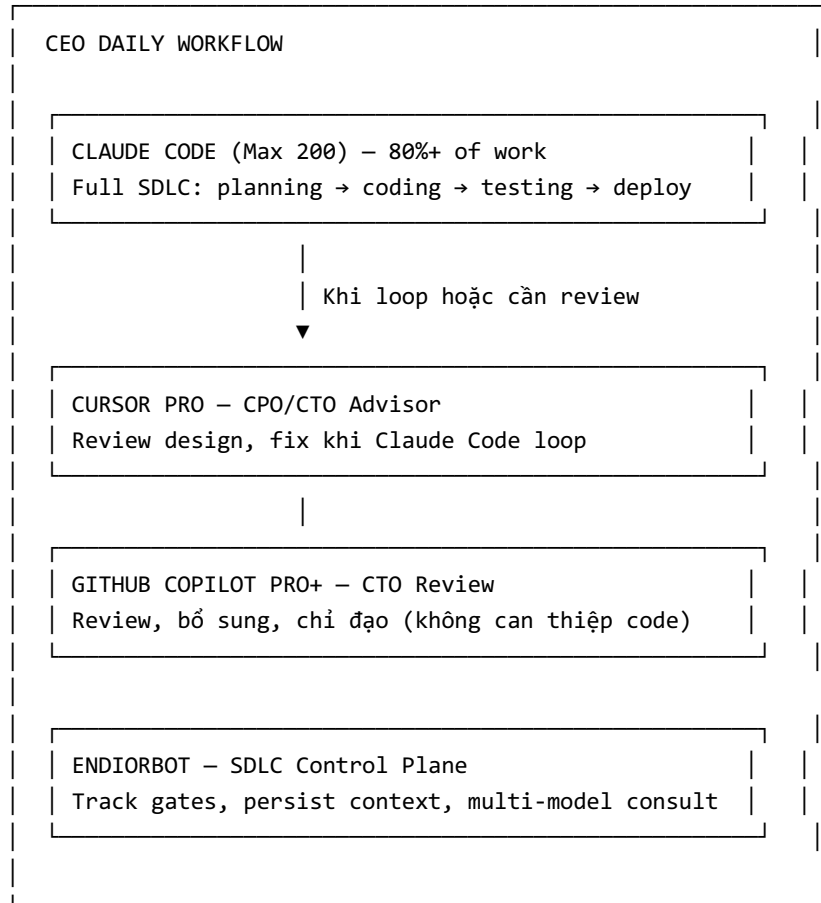
### IDENTITY (LOCKED - 1 SENTENCE)

EndiorBot là SDLC Control Plane cho Claude Code workflow, giúp CEO track gates + persist context + consult multi-model cho architecture decisions.

| Attribute       | Value                                                |
|-----------------|------------------------------------------------------|
| Version         | 3.1 FINAL (Expert-Reviewed)                          |
| Date            | February 28, 2026                                    |
| Status          | READY FOR EXECUTION                                  |
| Primary Tool    | Claude Code Max 200                                  |
| Secondary Tools | Cursor Pro (CPO/CTO), GitHub Copilot Pro+ (CTO)      |
| EndiorBot Role  | SDLC Control Plane + Brain + Multi-Model (arch only) |
| Runtime         | Cloud-first (Max 200 quota), Ollama fallback only    |

# 1. CEO Workflow (Reality)

EndiorBot phải FIT vào workflow hiện tại của CEO, không thay đổi nó:



## ⚠ FIT PRINCIPLE

EndiorBot KHÔNG thay thế Claude Code/Cursor/Copilot. EndiorBot BỔ SUNG cho workflow hiện tại với SDLC tracking + context persistence + architecture consultation.

## 2. Priority Stack (CEO Ranked)

| Priority | Capability                | Phase   | Scope                                                     |
|----------|---------------------------|---------|-----------------------------------------------------------|
| #1       | SDLC Control Plane        | Phase 1 | Track stages, gates, evidence across Claude Code sessions |
| #2       | Brain/Context Persistence | Phase 1 | Nhớ context giữa các sessions, inject vào new sessions    |
| #3       | Multi-Model Consultation  | Phase 2 | Claude + GPT + Gemini cho architecture decisions ONLY     |
| #4       | Handoff Orchestrator      | Later   | Route Claude Code → Cursor → Copilot khi loop (DEFER)     |

### Phase 1 Focus (Sprint 50-52)

- SDLC Control Plane: Gate checklist, evidence pointers, stage tracking
- Brain: Context persistence between Claude Code sessions
- NO: Multi-Model, NO: Team Presets, NO: SASE generation, NO: Vibecoding

### Phase 2 Focus (Sprint 53-55)

- Multi-Model Consultation: 2-model default (Claude lead + 1 expert)
- Architecture decisions only, NOT for coding tasks
- 4-model chỉ khi explicit flag --high-stakes

### Later (Defer)

- Handoff Orchestrator: CEO đã có manual workflow tốt
- Team Presets: Không cần 12 agents, CEO dùng Claude Code trực tiếp
- SASE Generation: Nice-to-have, không critical
- Desktop Shell: CLI đủ rồi

### 3. Three Invariants (LOCKED)

#### INVARIANT #1: Brain $\neq$ Truth

Brain chỉ ASSIST, không phải Source of Truth. SSOT nằm ở Git (commits, tags) và SDLC Framework (.sdlc-config.json). Brain có thể sai, SSOT không sai.

#### INVARIANT #2: Agent $\neq$ Authority

EndiorBot RECOMMEND, không APPROVE. Gate status là suggestion, CEO confirm. Không có 'gate pass' command — chỉ có 'gate recommend' (read-only). Authority = Human CEO.

#### INVARIANT #3: Approval = Human Confirm

Mọi action thay đổi state (gate pass, ADR approve, deploy trigger) đều yêu cầu explicit human confirmation. Agent chuẩn bị evidence + recommendation, CEO bấm confirm.

#### Derived Rules

- CLI command 'gate pass' → RENAMED to 'gate recommend'
- Gate approval → Git tag + evidence snapshot (CEO action)
- ADR approval → CEO confirm qua CLI prompt hoặc magic link
- No auto-approve, no auto-commit, no auto-deploy

## 4. SDLC Control Plane (#1 Priority)

Core capability: Track SDLC stages, gates, evidence across Claude Code sessions.

### 4.1 Stage Tracking

| Stage         | Tracking                  | Evidence        |
|---------------|---------------------------|-----------------|
| 00 FOUNDATION | Problem statement exists? | Link to doc     |
| 01 PLANNING   | Requirements documented?  | Link to specs   |
| 02 DESIGN     | Architecture decided?     | ADR links       |
| 03 INTEGRATE  | API contracts defined?    | Contract files  |
| 04 BUILD      | Implementation complete?  | PR links        |
| 05 TEST       | Tests passing?            | Coverage report |
| 06 DEPLOY     | Deployed?                 | Deploy log      |
| 07-09         | Deferred                  | Later phases    |

### 4.2 Gate Checklist (Read-Only)

endiorbot gate status G2

|                                                            |
|------------------------------------------------------------|
| GATE G2 (Design) – Project: BFlow                          |
| Checklist:                                                 |
| [✓] Architecture documented (ADR-045, ADR-046)             |
| [✓] Data model defined (docs/data-model.md)                |
| [X] API contracts finalized (missing: payment-api.yaml)    |
| [✓] Security review completed (security-review-2026-02.md) |
| Status: 3/4 complete                                       |
| Recommendation: READY with 1 exception                     |
| To approve: CEO must create Git tag 'G2-approved-bflow'    |

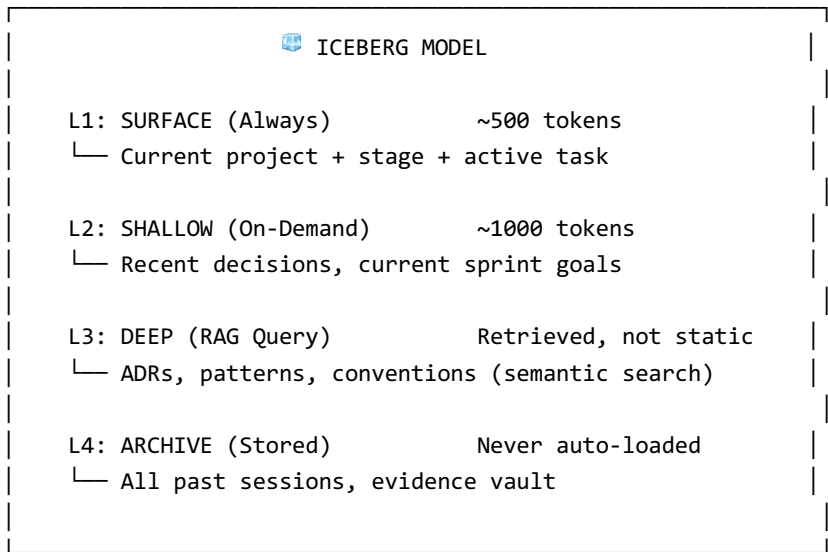
#### ⚠ NO 'gate pass' COMMAND

EndiorBot shows status + recommendation. CEO creates Git tag manually or via 'endiorbot gate confirm G2' (requires explicit --confirm flag).

## 5. Brain/Context Persistence (#2 Priority)

Core problem: Claude Code sessions không share context. CEO phải repeat context mỗi session.

### 5.1 4-Layer Iceberg (Simplified)



### 5.2 Key Changes from v3.0

| Aspect         | v3.0 (Wrong)         | v3.1 (Fixed)                 |
|----------------|----------------------|------------------------------|
| L3 injection   | Static ~3K tokens    | RAG query, semantic search   |
| Token budget   | \$10/day, ~200 turns | Per-session budget           |
| Context source | Brain as truth       | Brain assists, Git is truth  |
| History        | Inject all           | History Compactor summarizes |

### 5.3 Context Injection Flow

CEO: endiorbot context inject

- Load L1 (always): project ID, stage, active task
- Check task type:
  - Coding task → inject coding conventions (L2)
  - Architecture task → RAG query ADRs (L3)
  - Bug fix → RAG query recent issues (L3)
- Generate context block for Claude Code

4. Copy to clipboard / pipe to Claude Code

## 6. Multi-Model Consultation (#3 Priority)

### SCOPE LIMIT

Multi-Model Consultation ONLY for architecture decisions. NOT for coding, NOT for debugging, NOT for daily tasks. CEO đã có Claude Code Max 200 cho everyday work.

### 6.1 Model Configuration

| Model           | Role             | When Used                    |
|-----------------|------------------|------------------------------|
| Claude Opus 4.5 | Lead + Synthesis | Always (via Max 200)         |
| GPT-4o          | Expert #1        | Default 2nd model            |
| Gemini 2.0 Pro  | Expert #2        | Only with --high-stakes flag |
| Mistral Large   | Expert #3        | Only with --high-stakes flag |

### 6.2 Default: 2-Model (Cost-Efficient)

endiorbot consult "design payment gateway architecture"

- Queries: Claude Opus (lead) + GPT-4o (expert)
- Returns: Consolidated recommendation
- Cost: ~\$0.50-1.00 per consult

### 6.3 High-Stakes: 4-Model (Rare)

endiorbot consult --high-stakes "migrate to microservices?"

- Queries: Claude + GPT + Gemini + Mistral (parallel)
- Returns: Full consensus/disagreement analysis
- Cost: ~\$2-3 per consult
- Use case: Major architecture decisions, irreversible changes

### 💰 COST GATE

If consult would exceed \$5, prompt CEO for confirmation. Auto-downgrade to 2-model if CEO declines.



## 7. CLI Commands (Minimal)

### SDLC Control Plane

```
endiorbot stage status          # Current stage + checklist
endiorbot gate status G2       # Gate checklist (read-only)
endiorbot gate recommend G2     # Generate recommendation
endiorbot gate confirm G2 --confirm # CEO confirms (creates Git tag)
endiorbot evidence list        # List evidence pointers
endiorbot evidence add <link>  # Add evidence link
```

### Brain/Context

```
endiorbot context inject      # Generate context for Claude Code
endiorbot context status      # Show L1-L4 status
endiorbot context search "<query>" # RAG search L3/L4
endiorbot switch <project>    # Switch project context
```

### Multi-Model (Phase 2)

```
endiorbot consult "<question>"    # 2-model (Claude + GPT)
endiorbot consult --high-stakes    # 4-model parallel
```

### NOT in Phase 1

```
# DEFERRED:
# endiorbot team <preset>      # Team presets - LATER
# endiorbot vibe               # Vibecoding Index - LATER
# endiorbot mrp                # MRP generation - LATER
# endiorbot handoff            # Handoff orchestrator - LATER
```

## 8. Architecture (Cloud-First)

### 8.1 Runtime Model

| Component          | Runtime               | Notes                                  |
|--------------------|-----------------------|----------------------------------------|
| EndiorBot CLI      | Node.js (local)       | TypeScript, runs on CEO machine        |
| Claude API         | Cloud (Max 200)       | Primary AI, unlimited via subscription |
| GPT/Gemini/Mistral | Cloud (API)           | Multi-Model consultation               |
| Ollama             | Local (fallback)      | Only when cloud unavailable            |
| Brain Storage      | Local (~/.endiorbot/) | SQLite + JSON files                    |
| Evidence           | Git + Cloud links     | Pointers, not copies                   |

### 8.2 File Structure

```
~/.endiorbot/
├─ config.json      # API keys, defaults
├─ active.json      # Current project state (SSOT)
├─ brain/
│  └─ 11-surface.json # Always-loaded context
│  └─ 12-shallow.json # Recent decisions
│  └─ 13-index.db     # RAG index (SQLite + embeddings)
│  └─ 14-archive/     # Historical sessions
├─ projects/
│  └─ bflow/
│     └─ gates.json   # Gate status per project
│     └─ evidence.json # Evidence pointers
│  └─ nqh-bot/
└─ logs/             # Audit trail
```

### 8.3 active.json (Write-Lock)

```
{
  "activeProjectId": "bflow",
  "currentStage": "02-DESIGN",
  "lastCommandSource": "endiorbot-cli",
  "timestamp": "2026-02-28T16:50:00Z",
  "lock": {
    "isLocked": false,
    "lockedBy": null,
    "lockedAt": null
  }
}
```



## 9. Roadmap (Focused)

### PHASE 1 ONLY

Sprint 50-52 chỉ focus #1 SDLC Control Plane + #2 Brain. Không làm Multi-Model, không làm Team Presets, không làm SASE. Chứng minh value trước khi expand.

### Phase 1: Control Plane + Brain (Sprint 50-52)

| Sprint | Focus              | Deliverable                    | Success Metric              |
|--------|--------------------|--------------------------------|-----------------------------|
| 50     | SDLC Control Plane | stage status, gate recommend   | CEO tracks 1 project        |
| 51     | Brain L1-L2        | context inject, switch project | Context persists 7 days     |
| 52     | Brain L3 RAG       | context search, evidence list  | RAG retrieves relevant ADRs |

### Phase 2: Multi-Model (Sprint 53-55)

| Sprint | Focus                 | Deliverable            | Success Metric   |
|--------|-----------------------|------------------------|------------------|
| 53     | 2-Model Consultation  | consult command        | Response < 30s   |
| 54     | 4-Model (high-stakes) | --high-stakes flag     | Consensus report |
| 55     | ADR Generation        | ADR draft from consult | 80% usable draft |

### Later (DEFER)

- Team Presets: CEO không cần 12 agents, dùng Claude Code trực tiếp
- Handoff Orchestrator: Manual workflow đủ tốt
- SASE Generation: Nice-to-have
- Vibecoding Index: Not priority
- Desktop Shell: CLI đủ rồi

## 10. Success Metrics

### Phase 1 Complete (Sprint 52)

- CEO tracks at least 1 project through EndiorBot for 14 consecutive days
- Gate checklist accuracy: 80%+ items correct
- Context injection saves 5+ minutes per Claude Code session
- Brain persists context across sessions without manual re-entry

### Phase 2 Complete (Sprint 55)

- Multi-Model consultation: Response < 30 seconds
- 2-model default: Cost < \$1 per consult
- Architecture decision time: 30-60 min → 5 min
- ADR draft usability: 80% can be used with minor edits

### North Star (Weekly)

| Metric                    | Target        | Measurement            |
|---------------------------|---------------|------------------------|
| Hours saved/week          | 3+ hours      | CEO self-report        |
| Context injection usage   | Daily         | Command logs           |
| Gate tracking             | 100% projects | projects/ folder count |
| Multi-Model consultations | 2-3/week      | consult command count  |

# 11. Decisions Log

| Decision                   | Choice                             | Rationale                         |
|----------------------------|------------------------------------|-----------------------------------|
| Identity                   | SDLC Control Plane for Claude Code | Fit CEO workflow, not replace     |
| Priority #1                | SDLC tracking + gates              | CEO ranked highest                |
| Priority #2                | Brain/Context                      | Solve context persistence pain    |
| Priority #3                | Multi-Model                        | Only for architecture decisions   |
| Priority #4 (Defer)        | Handoff Orchestrator               | Manual workflow sufficient        |
| gate pass → gate recommend | Read-only by default               | Agent ≠ Authority invariant       |
| Multi-Model default        | 2 models (Claude + GPT)            | Cost-efficient, 4-model is opt-in |
| 12 Agents                  | DEFER                              | CEO uses Claude Code directly     |
| Team Presets               | DEFER                              | Not needed for CEO workflow       |
| Runtime                    | Cloud-first (Max 200)              | CEO already has subscription      |
| Ollama                     | Fallback only                      | Not primary runtime               |

 **MASTER PLAN v3.1 FINAL APPROVED**

EndiorBot = SDLC Control Plane cho Claude Code workflow. Focus Phase 1: Control Plane + Brain. Multi-Model chỉ Phase 2, chỉ cho architecture. Team Presets/Handoff/SASE: DEFER.

— End of Master Plan v3.1 FINAL —